



Apache Kafka Architectures and Fundamentals

Henrik Janzon, Solutions Engineer



Learning Objectives

After this module you will be able to:



- Give a high level description of the programming logic in Kafka producer and consumer clients



- Explain how EOS works to an interested lay person




- List the means with which Kafka provides durability and HA
- Illustrate on a high level, how you can secure your Kafka cluster

Apache Kafka is a Distributed Event Streaming Platform



Publish and subscribe to streams of events




Similar to a message queue or enterprise messaging system

Store streams of events



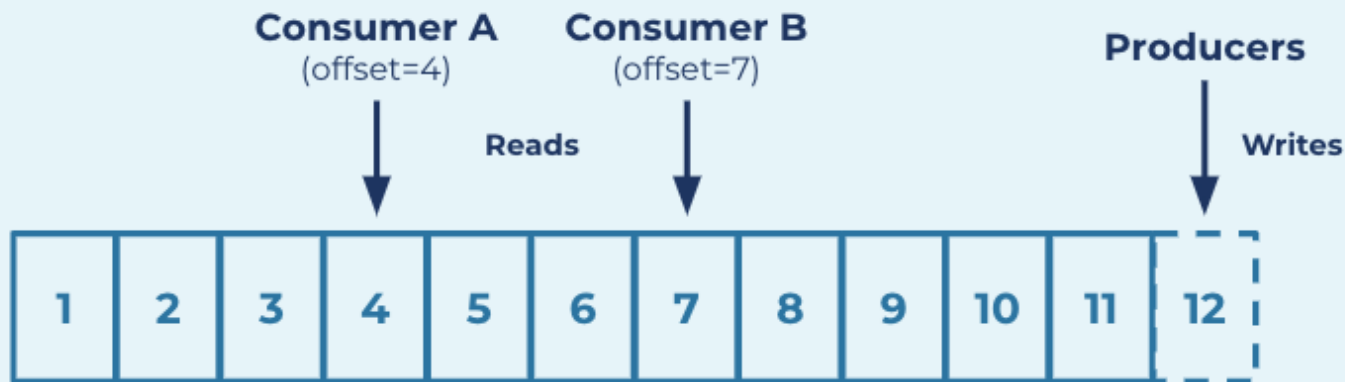

In a fault tolerant way

Process streams of events



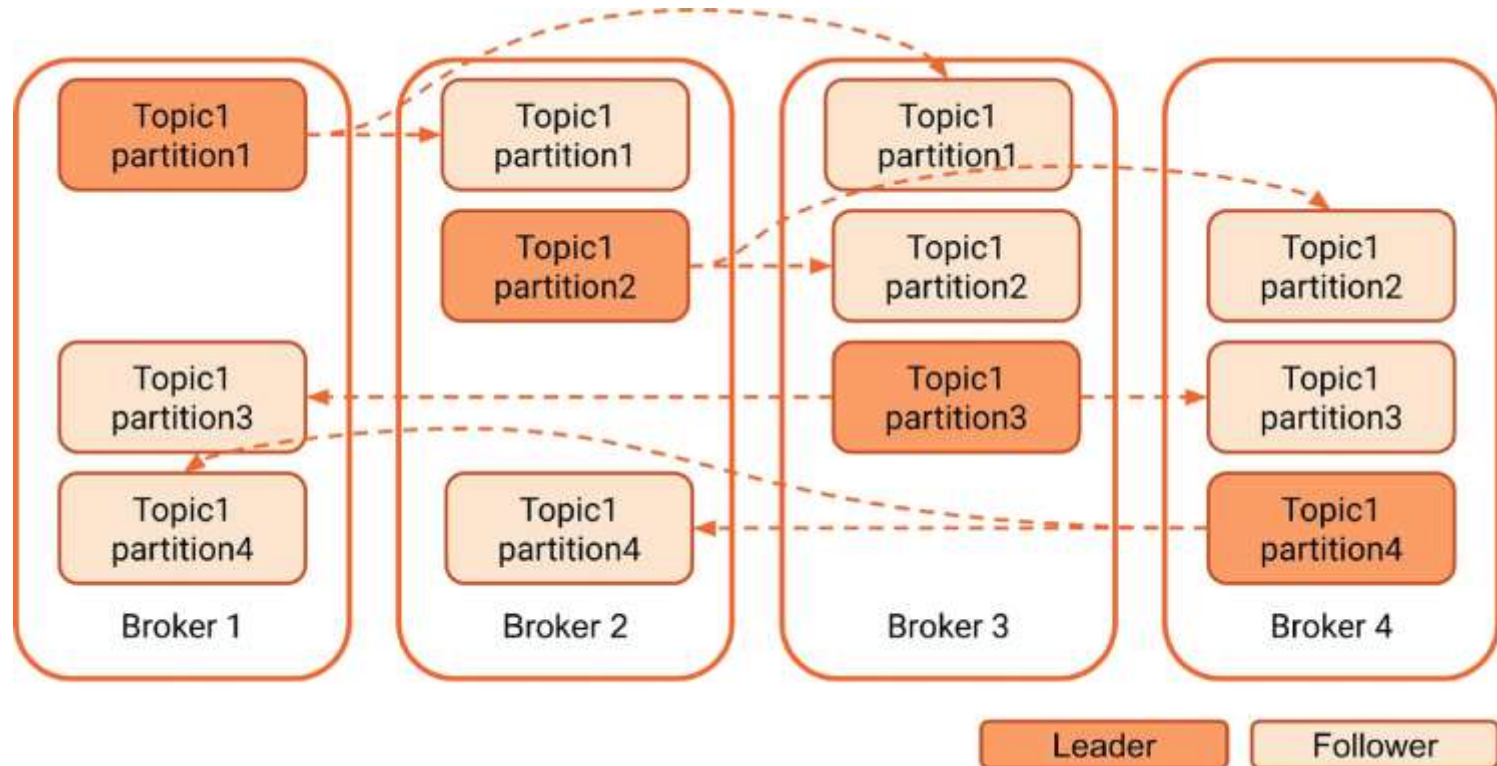
In real time, as they occur

Anatomy of a Kafka Topic



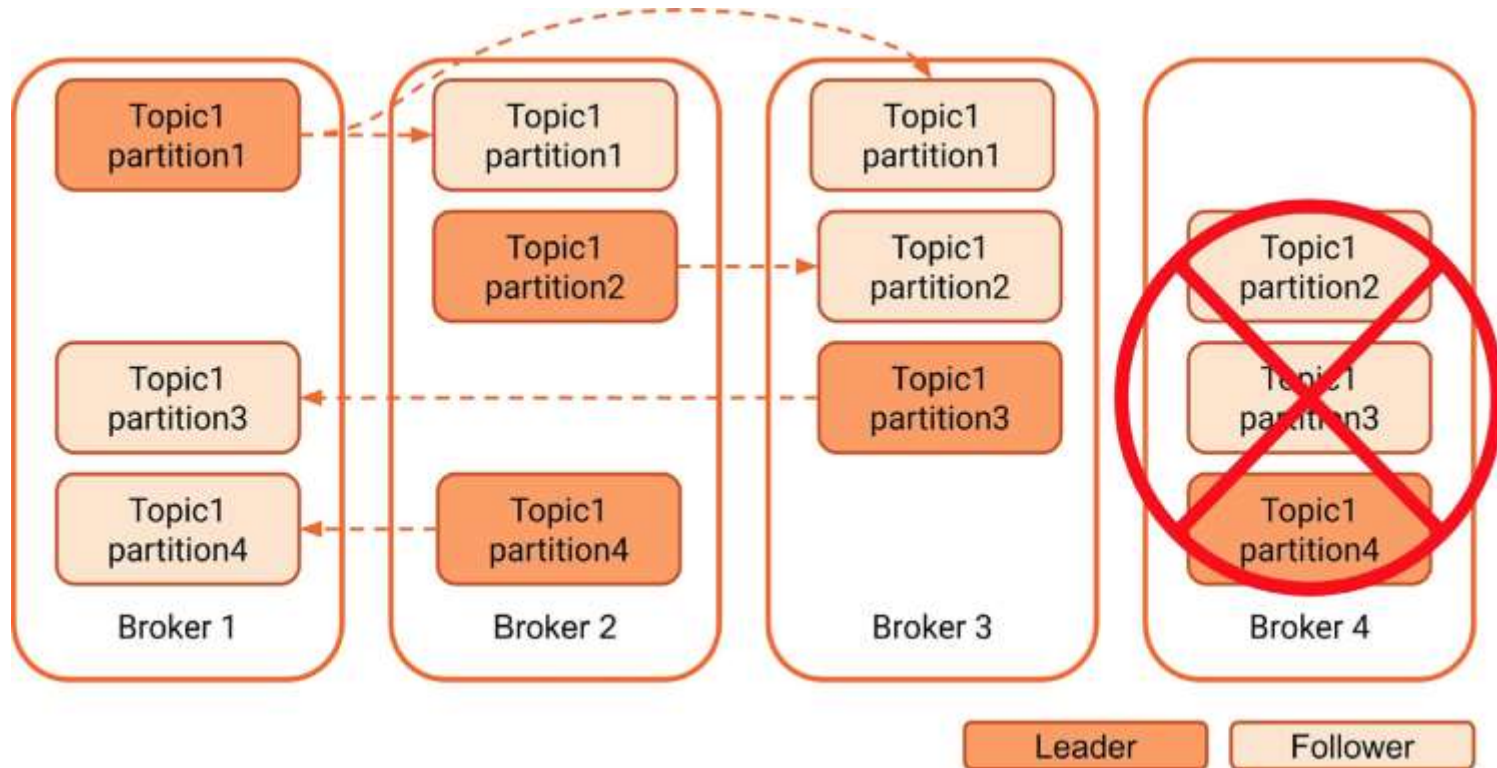


Partition Leadership & Replication





Partition Leadership & Replication

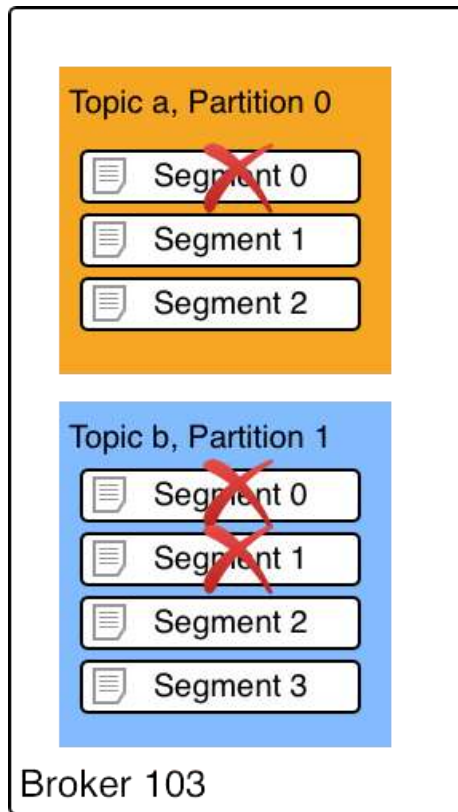




Data Retention Policy

How long do I want or can I store my data?

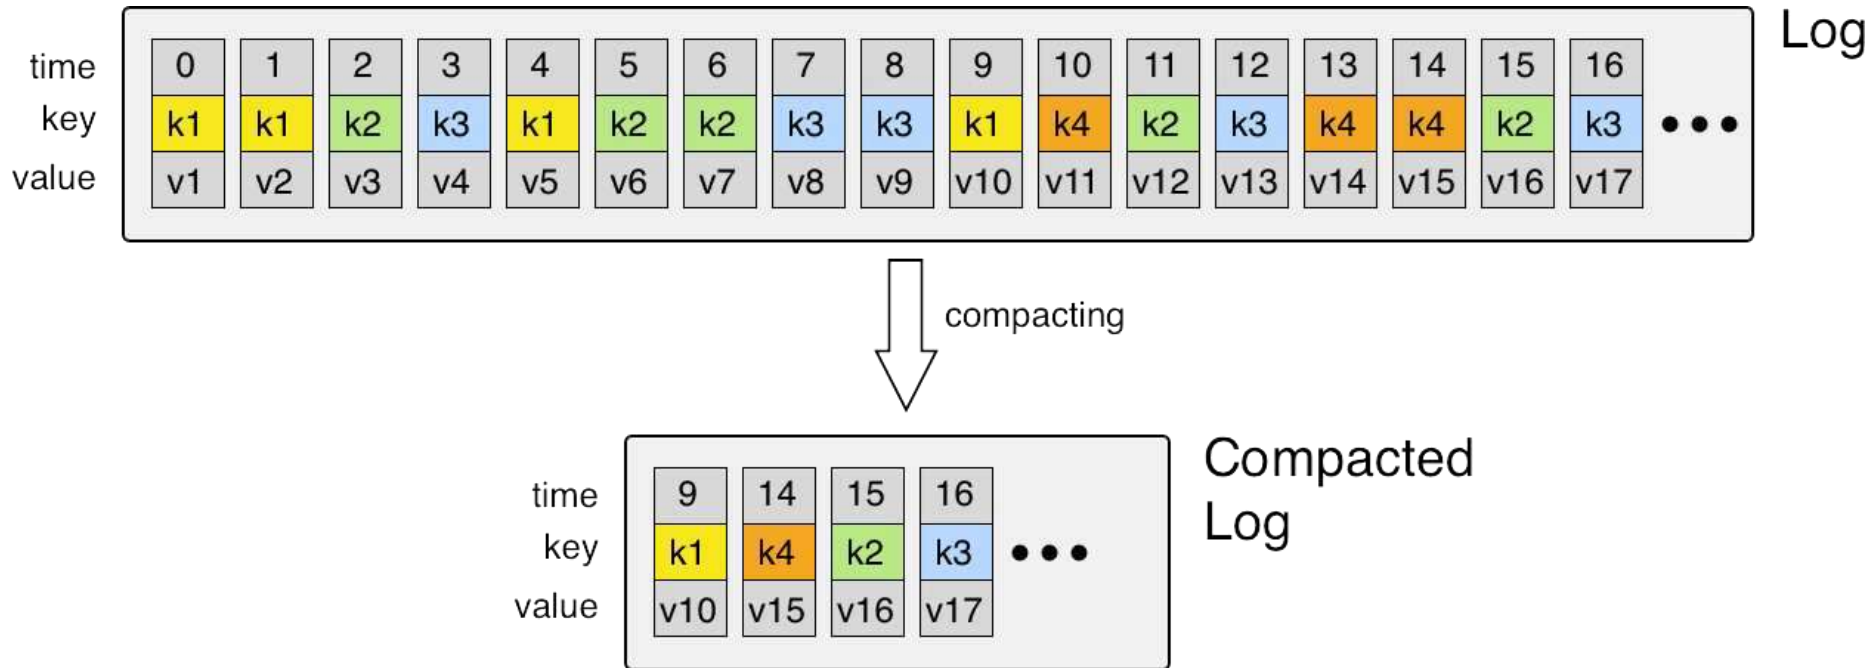
- How long (default: one week)
- Set **globally** or **per topic**
- Business decision
- Cost factor
- Compliance factor (e.g., GDPR)



Retention Policy



Compacted Topics



Development: A Basic Producer in Java

```
BasicProducer.java x
1 package clients;
2
3 import java.util.Properties;
4 import org.apache.kafka.clients.producer.KafkaProducer;
5 import org.apache.kafka.clients.producer.ProducerRecord;
6
7 public class BasicProducer {
8     public static void main(String[] args) {
9         System.out.println("*** Starting Basic Producer ***");
10
11         Properties settings = new Properties();
12         settings.put("client.id", "basic-producer-v0.1.0");
13         settings.put("bootstrap.servers", "kafka-1:9092,kafka-2:9092");
14         settings.put("key.serializer", "org.apache.kafka.common.serialization.StringSerializer");
15         settings.put("value.serializer", "org.apache.kafka.common.serialization.StringSerializer");
16
17         final KafkaProducer<String, String> producer = new KafkaProducer<>(settings);
18
19         Runtime.getRuntime().addShutdownHook(new Thread(() -> {
20             System.out.println("### Stopping Basic Producer ###");
21             producer.close();
22         }));
23
24         final String topic = "hello_world_topic";
25         for(int i=1; i<=5; i++){
26             final String key = "key-" + i;
27             final String value = "value-" + i;
28             final ProducerRecord<String, String> record = new ProducerRecord<>(topic, key, value);
29             producer.send(record);
30         }
31     }
32 }
```

configuration

create producer

shutdown behaviour

sending data



Development: A Basic Consumer in .NET/C#

```
8 namespace consumer_net {  
9     0 references  
    class Program {  
        0 references  
        static void Main (string[] args) {  
            Console.WriteLine ("Starting Consumer!");  
            var config = new Dictionary<string, object> {  
                { "group.id", "dotnet-consumer-group" },  
                { "bootstrap.servers", "kafka-1:9092" },  
                { "auto.commit.interval.ms", 5000 },  
                { "auto.offset.reset", "earliest" }  
            };  
  
            var deserializer = new StringDeserializer (Encoding.UTF8);  
            using (var consumer = new Consumer<string, string> (config, deserializer, deserializer)) {  
                consumer.OnMessage += (_, msg) =>  
                    Console.WriteLine ($"Read ('{msg.Key}', '{msg.Value}') from: {msg.TopicPartitionOffset}");  
  
                consumer.OnError += (_, error) =>  
                    Console.WriteLine ($"Error: {error}");  
  
                consumer.OnConsumeError += (_, msg) =>  
                    Console.WriteLine ($"Consume error ({msg.TopicPartitionOffset}): {msg.Error}");  
  
                consumer.Subscribe ("hello_world_topic");  
  
                while (true) {  
                    consumer.Poll (TimeSpan.FromMilliseconds (100));  
                }  
            }  
        }  
    }  
}
```

configuration

message handling

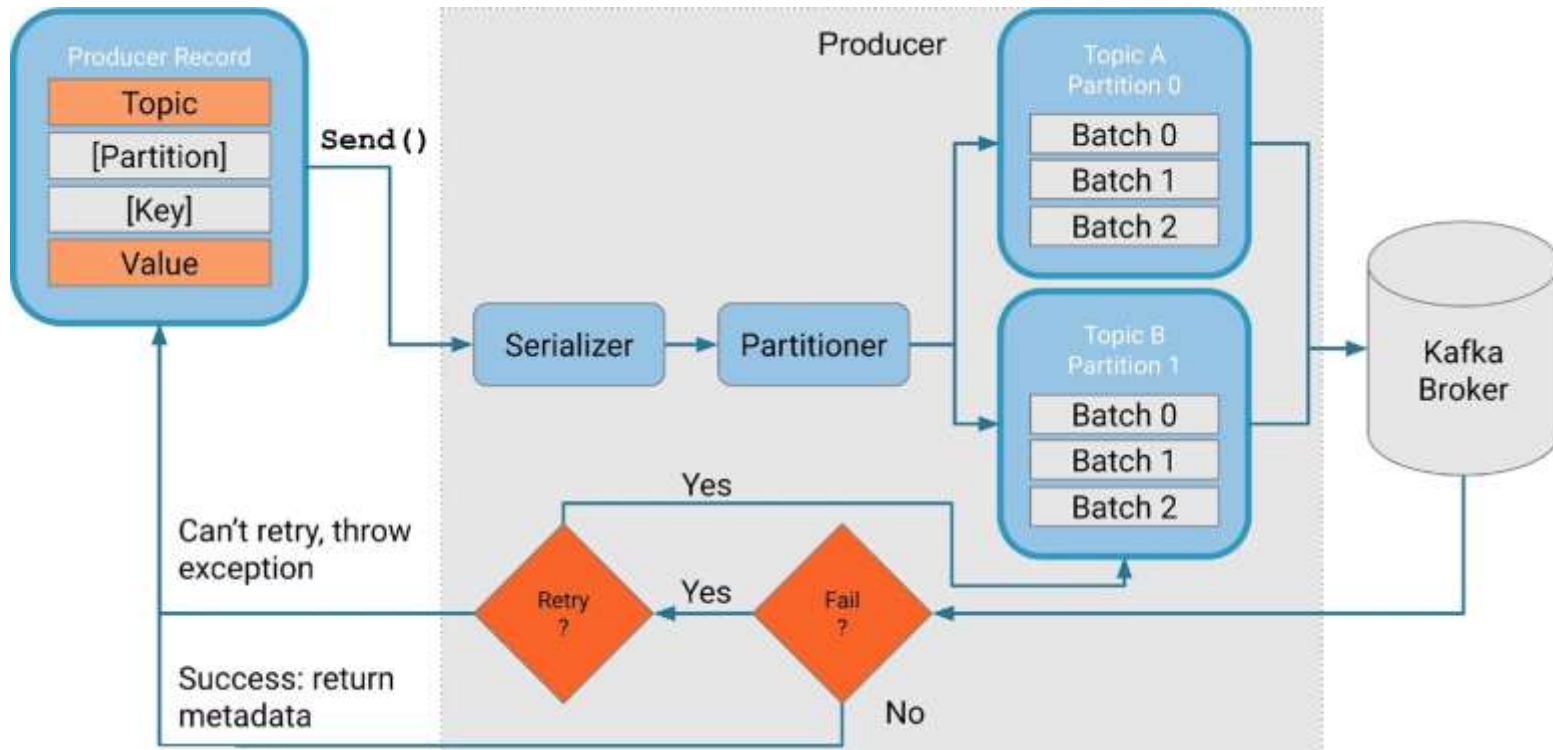
error handling

polling data

subscribing to topic




Producer Design







Producer Guarantees



Acks 0 (NONE)



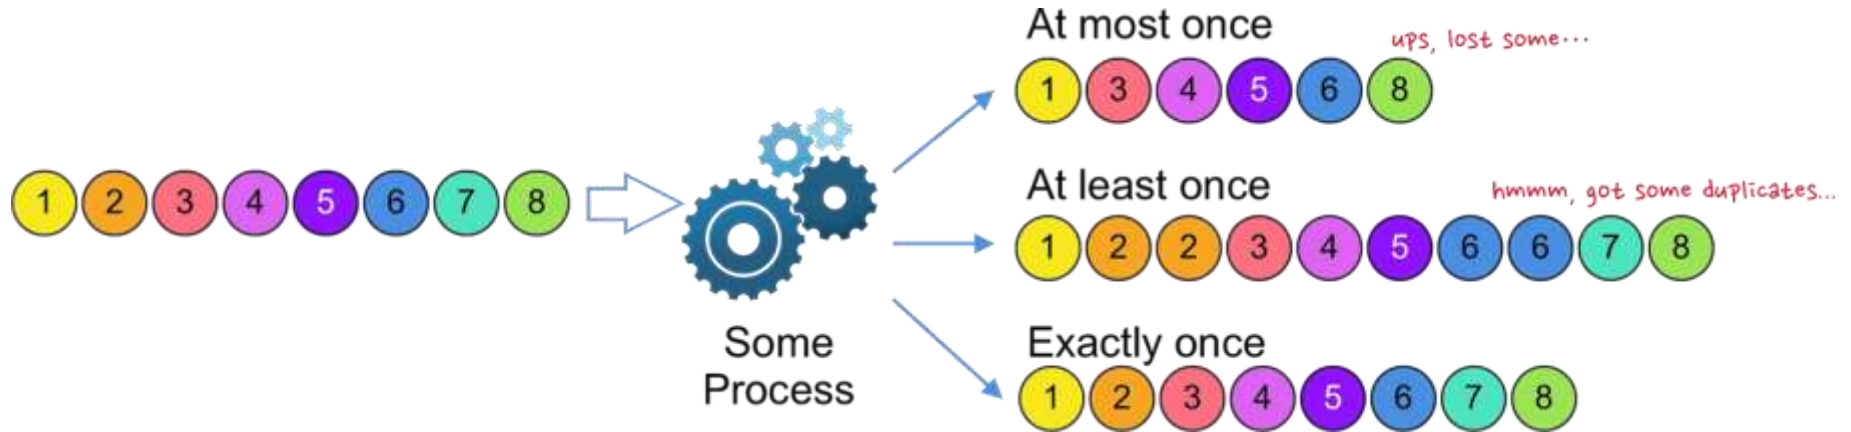
Acks 1 (LEADER)



Acks -1 (ALL)



Delivery Guarantees



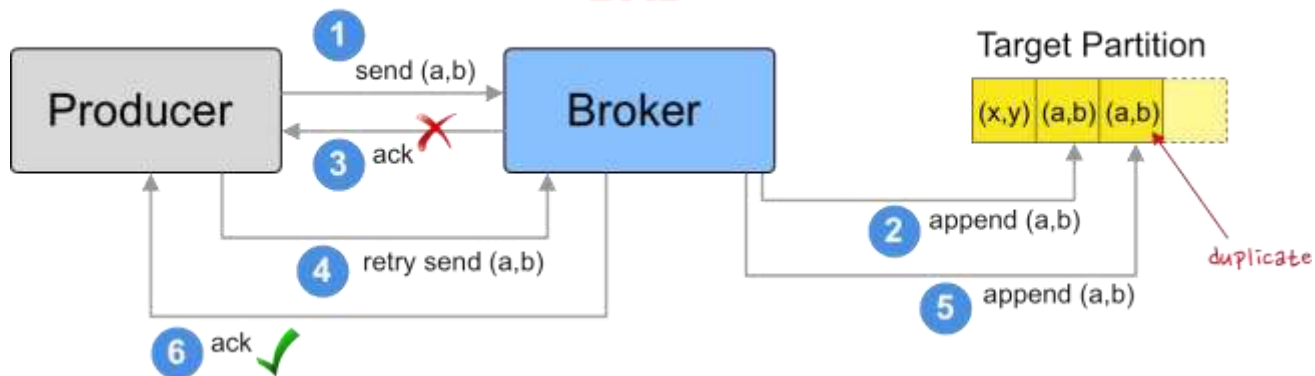
Idempotent Producers



GOOD



BAD





Exactly Once Semantics (EOS)

What is it?

- Strong **transactional guarantees** for Kafka
- Prevents clients from processing duplicate messages
- Handles failures gracefully

Use Cases


- Tracking ad views
- Processing financial transactions
- Stream processing



Partitioning Strategies

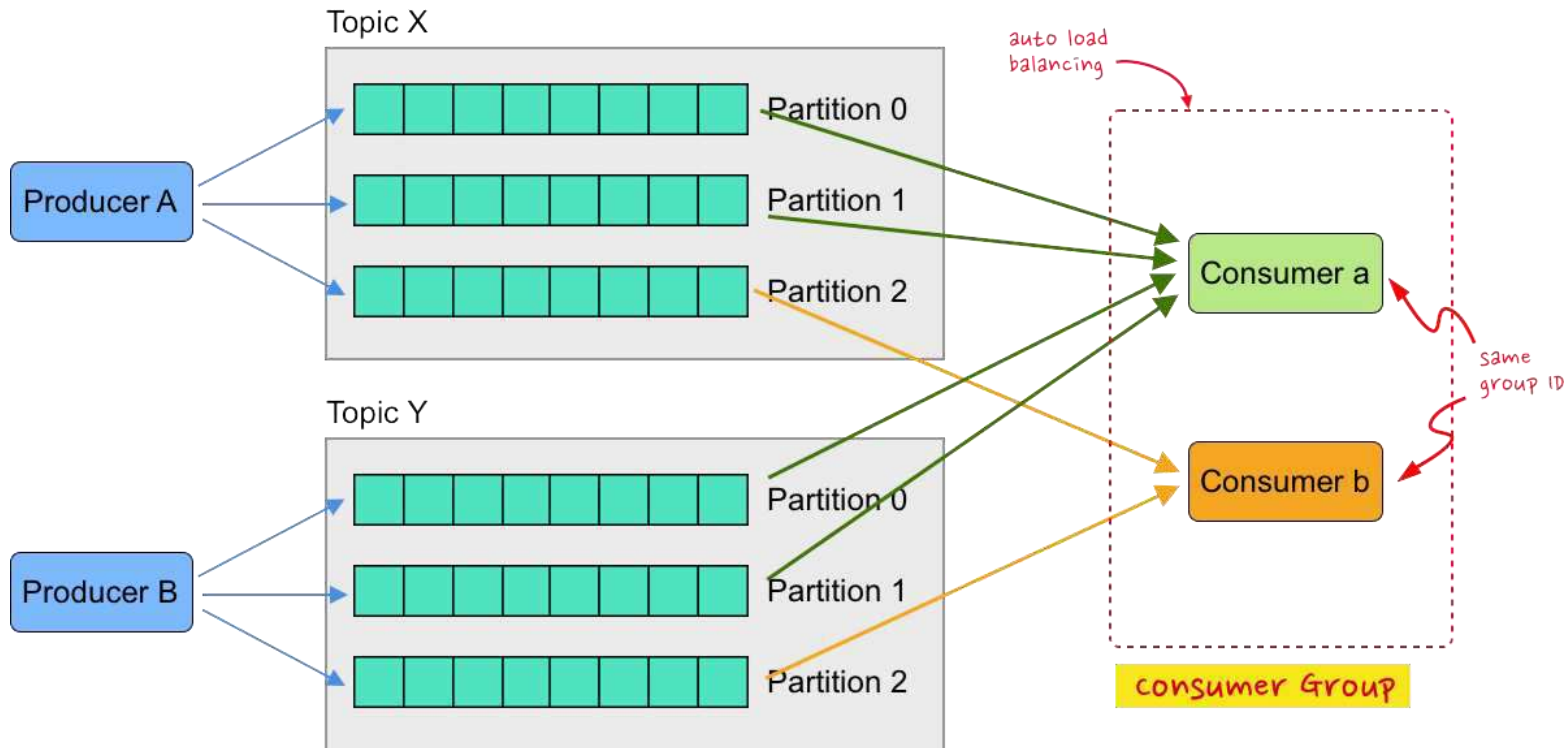
Why partitioning?

- Consumers need to **aggregate** or **join** by some key
- Consumers need **ordering guarantee**
- Concentrate data for **storage efficiency** and/or **indexing**








Consumer Groups



Consumer Rebalances





Security Overview

- Kafka supports Encryption in Transit
- Kafka supports Authorization and Authentication
- No Encryption at Rest out of the box
- Clients can be mixed with & without Encryption & Authentication



Client Side Security Features

- Encryption of Data in Transit
- Client Authentication
- Client Authorization



Encryption
in transit

SSL



authn & authz

authn: SASL or SSL

authz: ACLs



Keen to learn more?

Register for one of the Confluent Streaming Events

07 October - Middle East

12 October – Nordics

15 October – Rest of Europe

Visit: <https://events.confluent.io/>





Q&A



CONFLUENT